

# Filippos Christianos

Real-world  
Multi-Agent Systems:  
Research and Applications

## About me



Research Scientist at Huawei



PhD in the University of Edinburgh



Research Scientist, Intern at Nvidia



Electrical and Computer Engineering



Co-Author of the textbook “**Multi-Agent Reinforcement Learning**”  
[www.marl-book.com](http://www.marl-book.com) (MIT Press)

# In the next two hours...




**Part I** Introduction, Definitions, History of MARL

**Part II** Robotic Warehouses & Collaborative Environments:  
Scaling to Many Agents

**Part III** MARL in Practice

**Part IV** AI Agents: Introduction

**Part V** LLMs for Mobile-Phone Control

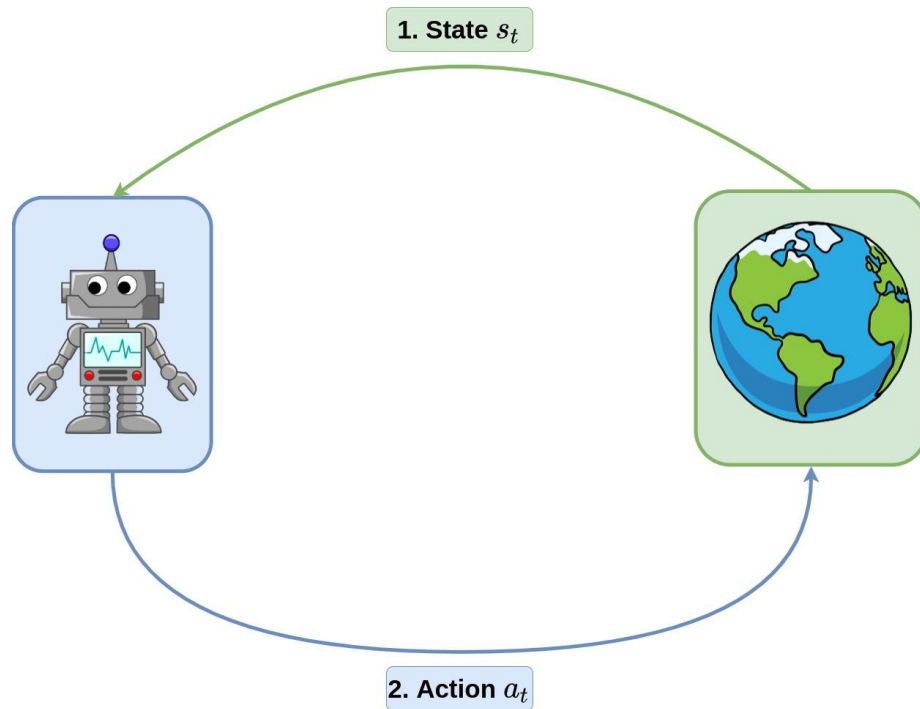


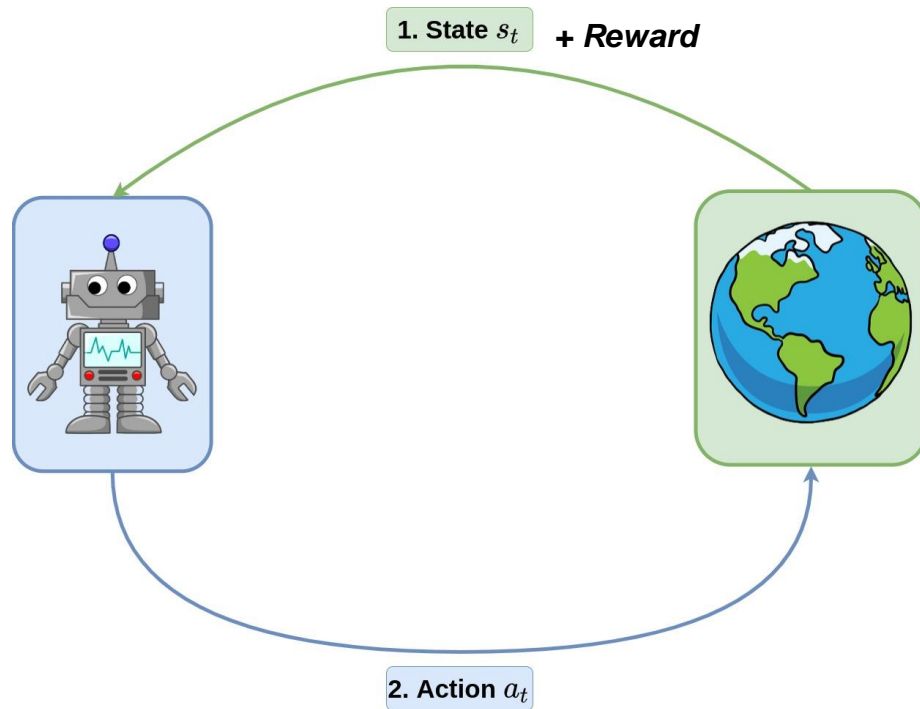
## Research Goal:

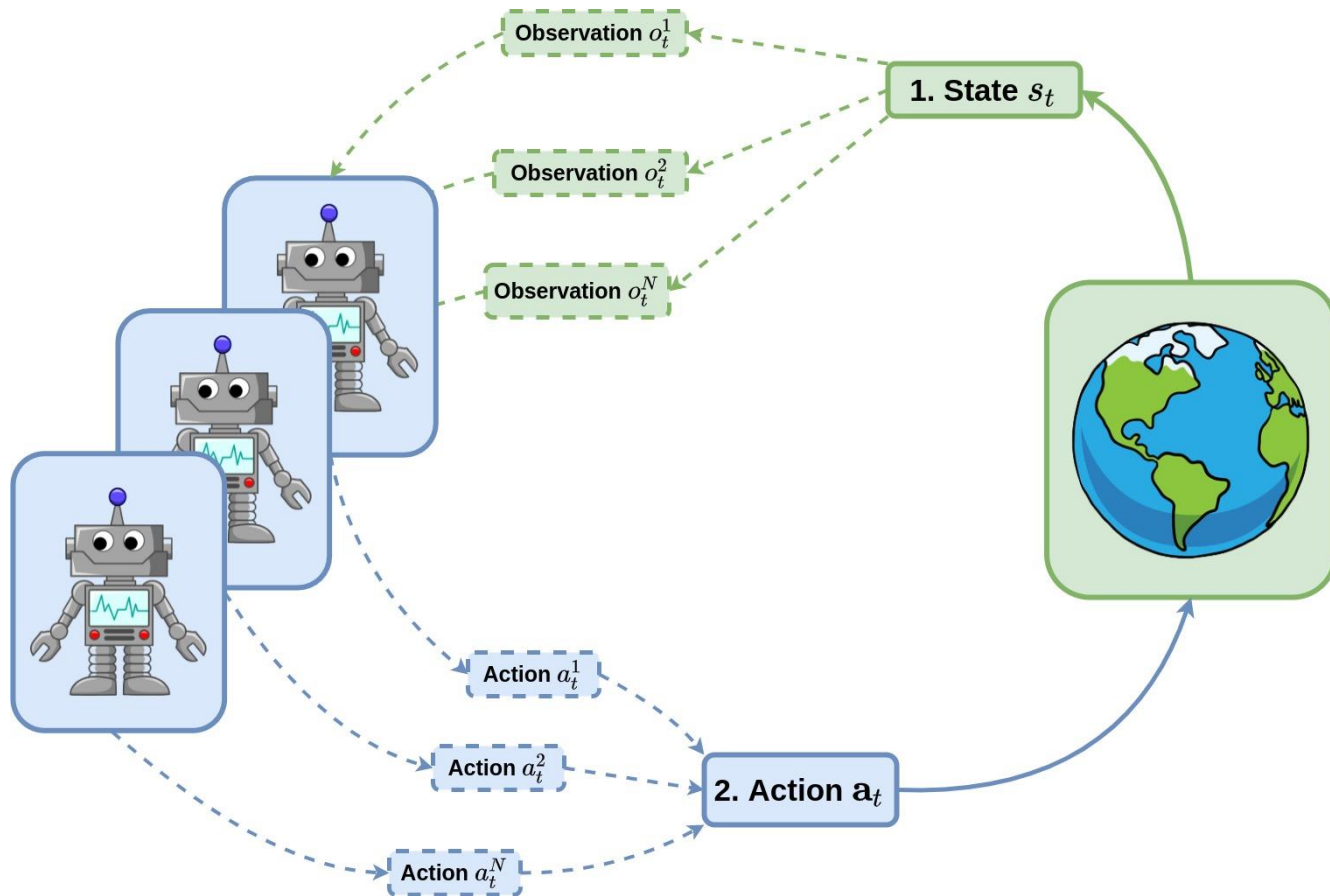
Create *autonomous agents* which can accomplish tasks in *complex dynamic environments*

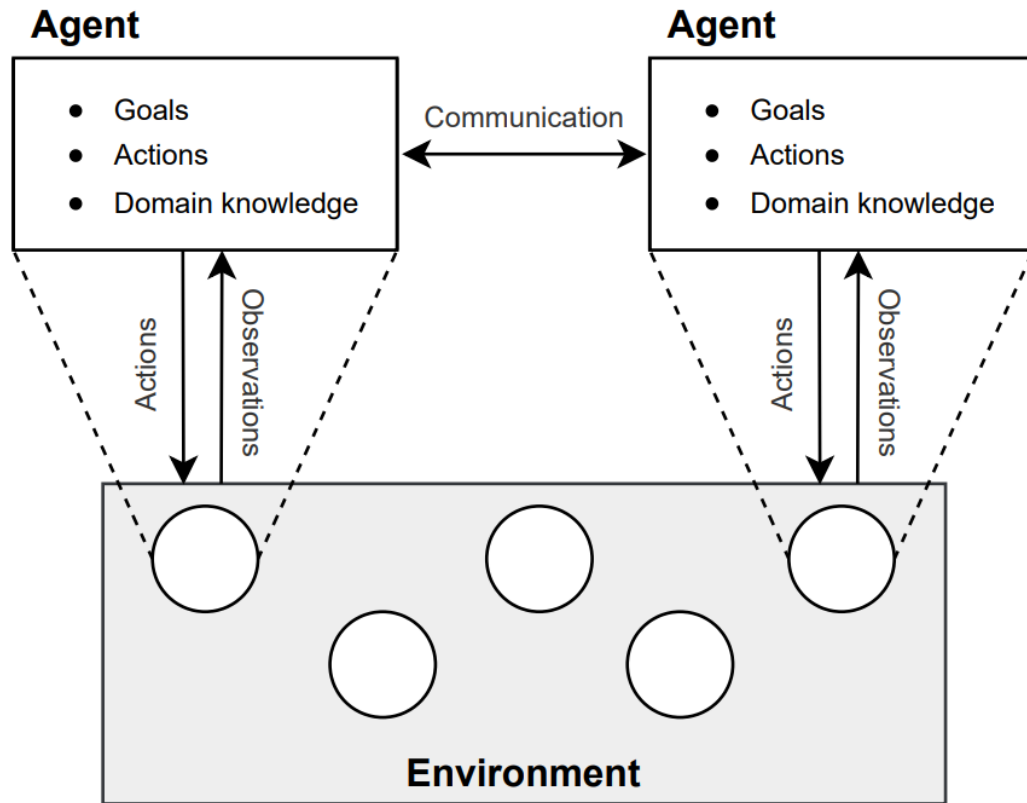
# Part I: Multi-Agent Systems – Scope and Definitions

1. What is an agent?
2. What is a multi-agent environment?
3. How do multiple agents *learn*?
4. Historical background and current state of research











# How did it start?

	R	P	S
R	0,0	-1,1	1,-1
P	1,-1	0,0	-1,1
S	-1,1	1,-1	0,0

**Rock Paper Scissors (RPS)**

	A	B
A	10	0
B	0	10

**Coordination Game**

	C	D
C	-1,-1	-5,0
D	0,-5	-3,-3

**Prisoner's Dilemma**

# Solution Concepts

Best Response

Nash Equilibrium

Minimax

Pareto Optimality

...No-Regret, Social Welfare, Fairness... *and more.*

# Solution Concepts

Best Response

Nash Equilibrium

Minimax

Pareto Optimality

...No-Regret, Social Welfare, Fairness... *and more.*

	R	P	S
R	0,0	-1,1	1,-1
P	1,-1	0,0	-1,1
S	-1,1	1,-1	0,0

Rock Paper Scissors (RPS)

# How did we get where we are now?


	R	P	S
R	0,0	-1,1	1,-1
P	1,-1	0,0	-1,1
S	-1,1	1,-1	0,0

Rock Paper Scissors (RPS)




# Two Types of Multi-agent Learning

**We control all agents:**



**We control one agent:**



# Two Types of Multi-agent Learning

## We control all agents:

How can agents learn to interact optimally?

- *Coordination and communication*
- *Efficient scaling to many agents*
- *Transfer abilities to other agents*



## We control one agent:

How can agent learn to collaborate with unknown other agents?

- *Recognising goals of other agents*
- *Integrating prediction and planning*



# Part II: Collaborative Training of Multiple Autonomous Agents

How can we train agents in large environments, and ensure they learn useful **joint policies**?





# What is the problem here?

- Many Agents!
  - The joint action space is huge
  - ... This makes the joint policy search space intractable

Can we do something?

# What is the problem here?

Can we do something?

Restricting the policy search  
space with parameter sharing

# What is the problem here?

Can we do something?

Restricting the policy search  
space with parameter sharing

Joint policy is constrained to consist of  
identical individual policies!

# Parameter Sharing

## No Parameter Sharing

Each agent has its own policy

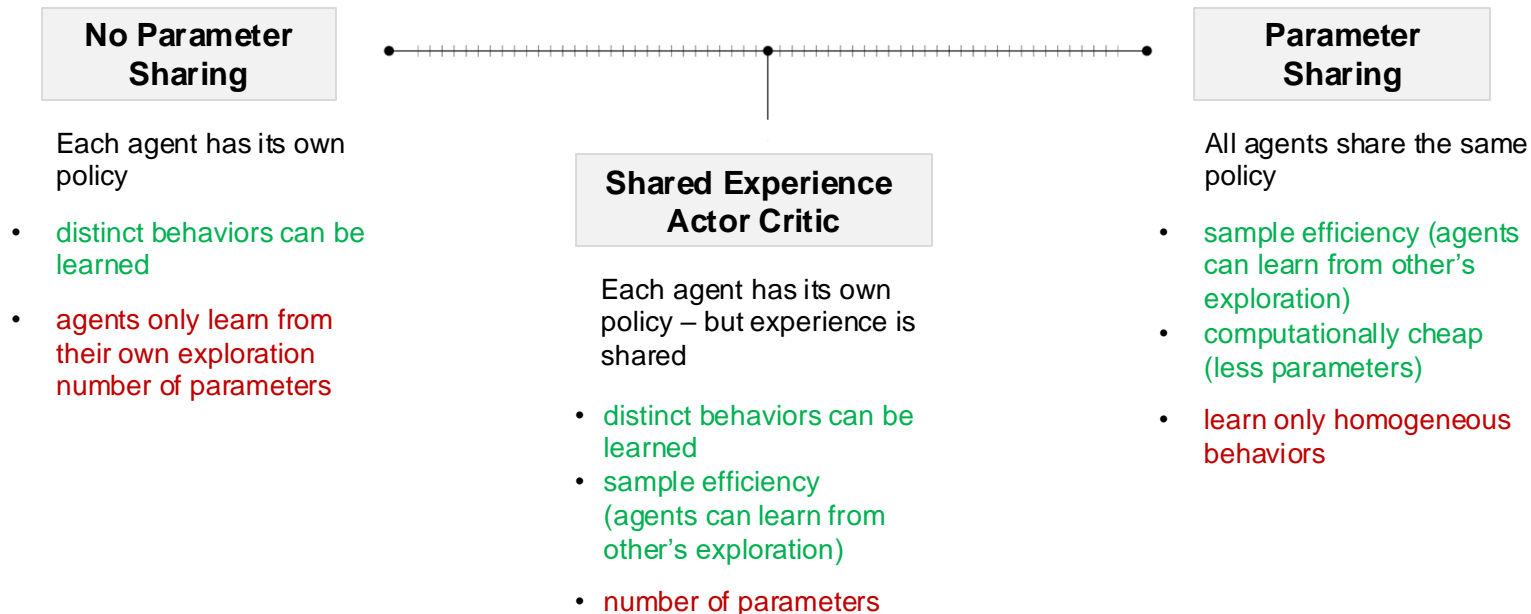
- distinct behaviors can be learned
- agents only learn from their own exploration  
number of parameters

## Parameter Sharing

All agents share the same policy


- sample efficiency (agents can learn from other's exploration)
- computationally cheap (less parameters)
- learn only homogeneous behaviors

# Parameter Sharing




# Motivational Example

- **Shared Experience Actor-Critic for Multi-Agent Reinforcement Learning**




# Motivational Example

- Both agents must reach goal simultaneously
- Sparse reward signal






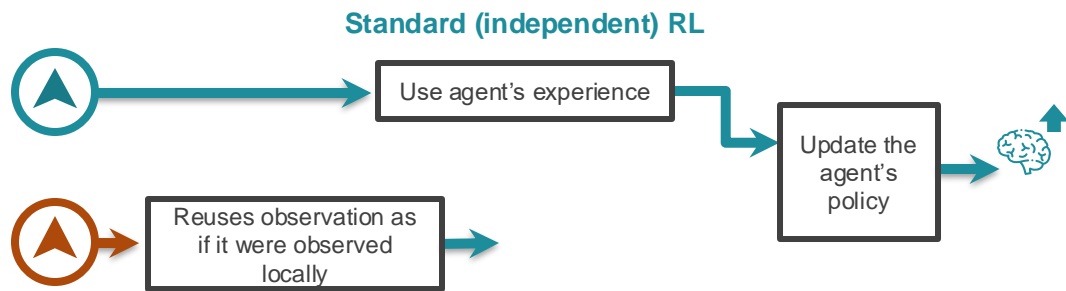
# Motivational Example


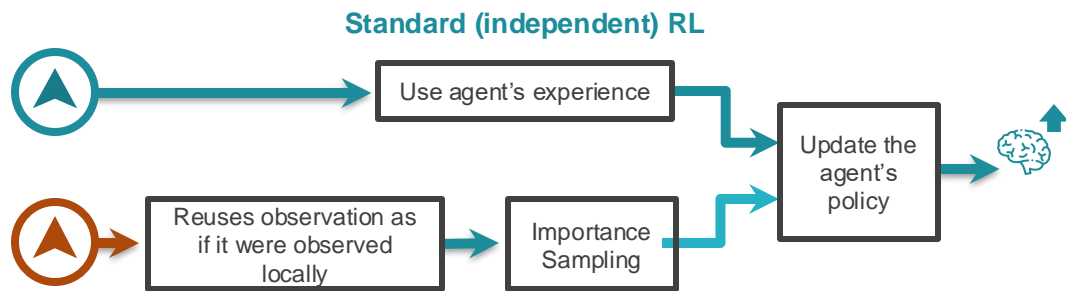
- Both agents must reach goal simultaneously
  - Sparse reward signal
- Idea: Make use of both agents' exploration
  - Share experience of agents


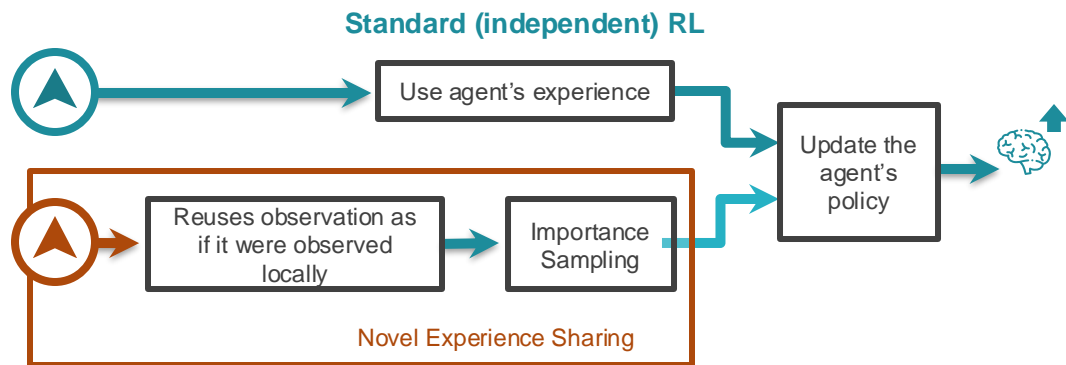













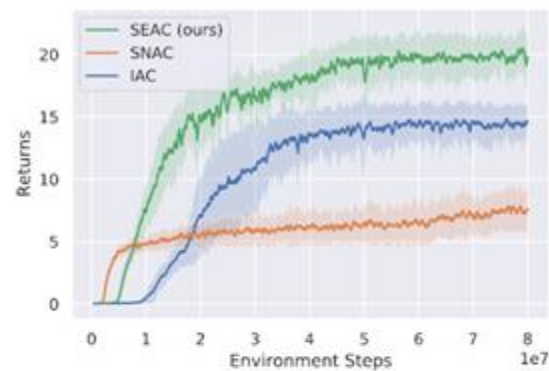


# A simulated robotics warehouse







(a) RWARE:  $(10 \times 11)$ , four agents




(b) RWARE:  $(10 \times 11)$ , two agents



(c) RWARE:  $(10 \times 11)$ , two agents, hard




(d) RWARE:  $(10 \times 20)$ , four agents




**Predator Prey (sparse)**



**SMAC – 3m (sparse)**



**Level-Based Foraging (LBF)**




**Multi-Robot Warehouse  
(RWARE)**

	IAC	SNAC	SEAC (ours)	QMIX	MADDPG	ROMA
PP (sparse)	-0.04 $\pm$ 0.13	-0.04 $\pm$ 0.1	<b>1.93 <math>\pm</math>0.13</b>	0.05 $\pm$ 0.07	<b>2.04 <math>\pm</math>0.08</b>	0.04 $\pm$ 0.07
SMAC-3m (sparse)	-0.13 $\pm$ 0.01	-0.14 $\pm$ 0.02	<b>-0.03 <math>\pm</math>0.03</b>	<b>0.00 <math>\pm</math>0.00</b>	<b>-0.01 <math>\pm</math>0.01</b>	<b>0.00 <math>\pm</math>0.00</b>
LBF-(15x15)-3ag-4f	0.13 $\pm$ 0.04	0.18 $\pm$ 0.08	<b>0.43 <math>\pm</math>0.09</b>	0.03 $\pm$ 0.01	0.01 $\pm$ 0.02	0.03 $\pm$ 0.02
LBF-(8x8)-2ag-2f-coop	0.37 $\pm$ 0.10	0.38 $\pm$ 0.10	<b>0.64 <math>\pm</math>0.08</b>	<b>0.79 <math>\pm</math>0.31</b>	0.01 $\pm$ 0.02	0.01 $\pm$ 0.02
RWARE-(10x20)-4ag	13.75 $\pm$ 1.26	9.53 $\pm$ 0.83	<b>23.96 <math>\pm</math>1.92</b>	0.00 $\pm$ 0.00	0.00 $\pm$ 0.00	0.00 $\pm$ 0.00
RWARE-(10x11)-4ag	<b>40.10 <math>\pm</math>5.60</b>	36.79 $\pm$ 2.36	<b>45.11 <math>\pm</math>2.90</b>	0.00 $\pm$ 0.00	0.00 $\pm$ 0.00	0.01 $\pm$ 0.01




# Analysis



**Best vs. Worst performing agents  
on RWARE, (10x20), four agents**


# Analysis



## Best vs. Worst performing agents on RWARE, (10x20), four agents

- Agents learn simultaneously which helps in exploring promising joint actions more

# Analysis



## Best vs. Worst performing agents on RWARE, (10x20), four agents

- Agents learn simultaneously which helps in exploring promising joint actions more
- Synchronise training progress of agents

# Environments with heterogeneous agents

### No Parameter Sharing

Each agent has its own policy

- distinct behaviors can be learned
- agents only learn from their own exploration  
number of parameters

### Shared Experience Actor Critic

Each agent has its own policy – but experience is shared

- distinct behaviors can be learned
- sample efficiency (agents can learn from other's exploration)
- number of parameters

### Parameter Sharing

All agents share the same policy

- sample efficiency (agents can learn from other's exploration)
- computationally cheap (less parameters)
- learn only homogeneous behaviors

## Selective Parameter Sharing

Groups of agents that share policies

- distinct behaviors can be learned
- sample efficiency (agents can learn from other's exploration)
- computationally cheap

## No Parameter Sharing

Each agent has its own policy

- distinct behaviors can be learned
- agents only learn from their own exploration  
number of parameters

## Shared Experience Actor Critic

Each agent has its own policy – but experience is shared

- distinct behaviors can be learned
- sample efficiency (agents can learn from other's exploration)
- number of parameters

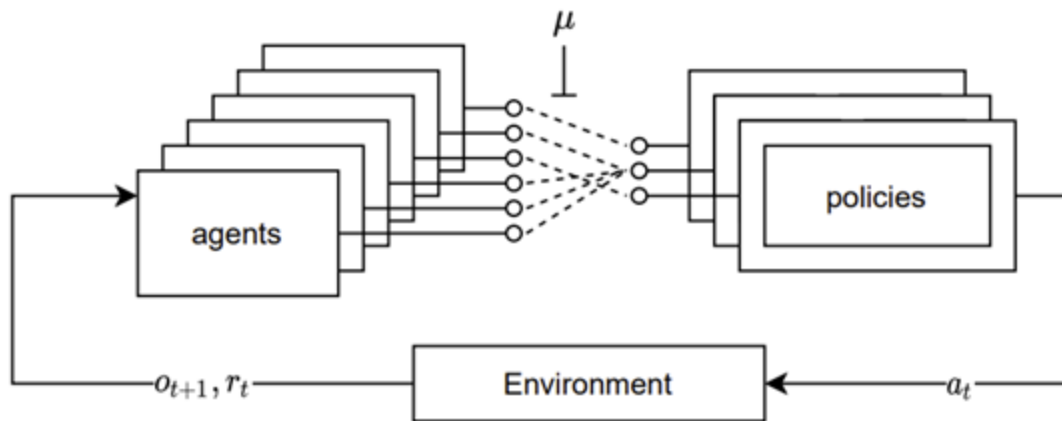
## Parameter Sharing

All agents share the same policy

- sample efficiency (agents can learn from other's exploration)
- computationally cheap (less parameters)
- learn only homogeneous behaviors


# Selective Parameter Sharing

But we can apply it selectively.




# Selective Parameter Sharing

We identify agents with similar reward and observation transition functions and have them share parameters.







# Visualising the embedding space




# Experiments: Environments



(a) Blind-Particle Spread



(b) Coloured Multi-Robot Warehouse




(c) Level-based Foraging




(d) SMAC


	# Agents	# Types	Type Distribution
BPS (1)	15	3	5-5-5
BPS (2)	30	3	10-10-10
BPS (3)	30	5	6-6-6-6-6
BPS (4)	30	5	2-2-2-15-9
BPS-h (1)	15	3 <sup>†</sup>	5-5-5
BPS-h (2)	30	5 <sup>†</sup>	6-6-6-6-6
BPS-h (3)	200	4 <sup>†</sup>	50-50-50-50
C-RWARE (1)	4	2 <sup>‡</sup>	2-2
C-RWARE (2)	8	2 <sup>‡</sup>	4-4
C-RWARE (3)	16	2 <sup>‡</sup>	8-8
LBF	12	3	4-4-4-4
MMM2	10	3 <sup>§</sup>	7-2-1




(a) BPS (3)




(b) BPS-h (2)




(c) C-RWARE (1)



(d) C-RWARE (3)



(e) LBF




(f) SMAC (MMM2)




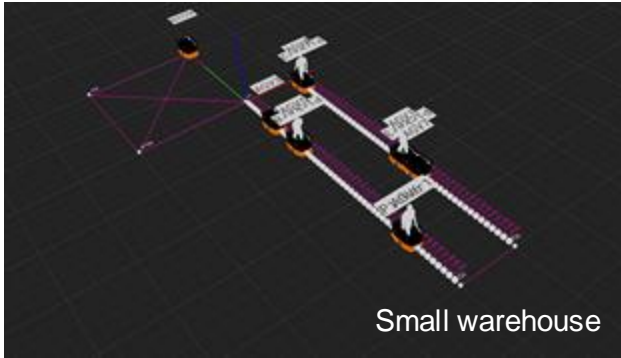
# Scalable Multi-Agent Reinforcement Learning for Warehouse Logistics with Robotic and Human Co-Workers



Aleksandar Krnjaic, Raul D. Steleac,  
Jonathan D. Thomas, Georgios Papoudakis,  
Lukas Schäfer, Andrew Wing Keung To,  
Kuan-Ho Lao, Murat Cubuktepe,  
Matthew Haley, Peter Börsting,  
Stefano V. Albrecht




# Warehouse Simulator



# Evaluation Results

1. Comparison against the two industry leading heuristics:
  - a. Follow Me (FM)
  - b. Pick Don't Move (PDM)
2. Ablation study of the hierarchical module for three data sharing mechanisms:
  - a. Independent Actor-Critic (IAC)
  - b. Shared Network Actor-Critic (SNAC)
  - c. Shared Experience Actor Critic (SEAC)



# Part III: Algorithms in Practice

Implementing a MARL  
algorithm in PyTorch



# Implementing MARL Algorithms

<https://github.com/marl-book/codebase>

```
import lbforaging
import gym

env = gym.make("Foraging-8x8-2p-1f-v2")
```

```
1 env.observation_space
2 >> Tuple(Box(..., 15), Box(..., 15))
3
4 env.action_space
5 >> Tuple(Discrete(6), Discrete(6))
```

```
1 observations = env.reset()
2 next_observations, rewards, terminal_signal, _ = env.step(
    actions)
```

# MARL: Neural Networks

```
import torch
from torch import nn
from typing import List

class MultiAgentFCNetwork(nn.Module):
    def __init__(
        self,
        in_sizes: List[int],
        out_sizes: List[int]
    ):
        super().__init__()


        # We use the ReLU activation function:
        activ = nn.ReLU()
        # We use two hidden layers of 64 units each:
        hidden_dims = (64, 64)

        n_agents = len(in_sizes)
        # The number of agents is the length of the
        # input and output vector
        assert n_agents == len(out_sizes)

        # We will create 'n_agents' (independent) networks
        self.networks = nn.ModuleList()


        # For each agent:
        for in_size, out_size in zip(in_sizes, out_sizes):
            network = [
                nn.Linear(in_size, hidden_dims[0]),
                activ(),
                nn.Linear(hidden_dims[0], hidden_dims[1]),
                activ(),
                nn.Linear(hidden_dims[1], out_size),
            ]
            self.networks.append(nn.Sequential(*network))

    def forward(self, inputs: List[torch.Tensor]):
        # The networks can run in parallel:
        futures = [
            torch.jit.fork(model, inputs[i])
            for i, model in enumerate(self.networks)
        ]
        results = [torch.jit.wait(fut) for fut in futures]
        return results
```



# Seamless Parameter Sharing Implementation

```
1 class MultiAgentFCNetwork_SharedParameters(nn.Module):
2
3     def __init__(
4         self,
5         in_sizes: List[int],
6         out_sizes: List[int]
7     ):
8
9         # ... same as MultiAgentFCNetwork
10
11         # We will create one (shared) network
12         # This assumes that input and output size of the
13         # networks is identical across agents. If not, one
14         # could first pad the inputs and outputs
15
16         network = [
17             # ... same as MultiAgentFCNetwork
18         ]
19         self.network = nn.Sequential(*network)
20
21     def forward(self, inputs: List[torch.Tensor]):
22
23         # A forward pass of the same network in parallel
24         futures = [
25             torch.jit.fork(self.network, inp)
26             for inp in inputs
27         ]
28         results = [torch.jit.wait(fut) for fut in futures]
29         return results
```



# Initialising and Querying the Models

```
1 # Example of observation of agent 1:
2 # obs1 = torch.tensor([1, 0, 2, 3, 0])
3
4 # Example of observation of agent 2:
5 # obs2 = torch.tensor([0, 0, 0, 3, 0])
6
7 obs_sizes = (5, 5)
8
9 # Example of action of agent 1:
10 # act1 = [0, 0, 1] # one-hot encoded
11
12 # Example of action of agent 2:
13 # act2 = [1, 0, 0] # one-hot encoded
14
15 action_sizes = (3, 3)
16
17 model = MultiAgentFCNetwork(obs_sizes, action_sizes)
18
19 # Alternatively, the shared parameter model can be used instead:
20 # model = MultiAgentFCNetwork_SharedParameters(
21 #     obs_sizes, action_sizes
22 # )
23
24
25 1 # obs1, obs2, model as above
26 2
27 3 q_values = model([obs1, obs2])
28 4 >> ([Q11, Q12, Q13], [Q21, Q22, Q23])
29 5 # where Qij is the Q value of agent i doing action j
```

# Independent DQN

```
1 # obs1, obs2, model as above
2
3 q_values = model([obs1, obs2])
4 >> ([Q11, Q12, Q13], [Q21, Q22, Q23])
5 # where Qij is the Q value of agent i doing action j
```


```
1 # we are creating a new "agent" dimension
2 q_values_stacked = torch.stack(q_values)
3 print(q_values_stacked.shape)
4 >> [2, 3]
5 # 2: agent dimension, 3: action dimension
6
7 # calculating best actions per agent (index):
8 _, a_prime = q_values_stacked.max(-1)
```

# Optimising...

```
1 params_list = list(nn_agent1.parameters())
2               + list(nn_agent2.parameters())
3               + list(nn_agent3.parameters())
4               + ...
5 common_optimiser = torch.optim.Adam(params_list)
6 ...
7 loss = loss1 + loss2 + loss3 + ...
8 loss.backward()
9 common_optimizer.step()
```

# Part IV: AI Agents

**Task:** "Create a model that predicts the optimal price for real estate properties based on features like location, size, and market trends."

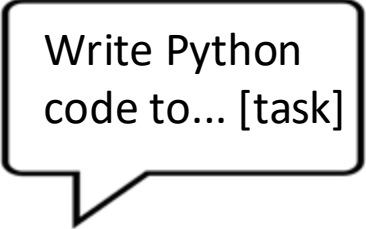




## We could have an LLM-based AI Agent

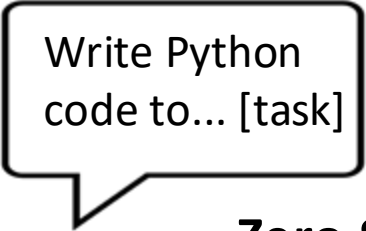
1. Make use of general LLMs as priors – they encode general enough priors
2. Allow agents to fix their internal structures for specific tasks

# LLM-based Agent



Write Python  
code to... [task]

# LLM-based Agent



Write Python  
code to... [task]

**Zero-Shot**

# LLM-based Agent

Write Python  
code to... [task]

**Zero-Shot**



START



FINISH


# LLM-based Agent

Write Python  
code to... [task]

**Zero-Shot**




START



FINISH

**Agentic Workflow**



What are structuring techniques that people have been using?

# Emerging components of an AI Agent

1. Planning
2. Reflection
3. Tools
4. Memory
5. Multi-Agent

# Planning



Recommended reading:

**Chain-of-Thought Prompting Elicits Reasoning in Large Language Models.** Wei et al.

**Self-Consistency Improves Chain of Thought Reasoning in Language Models.** Wang et al.



# Planning

Do a linear regression in NumPy in this dataset.



Recommended reading:

**Chain-of-Thought Prompting Elicits Reasoning in Large Language Models.** Wei et al.

**Self-Consistency Improves Chain of Thought Reasoning in Language Models.** Wang et al.

# Planning

Do a linear regression in NumPy in this dataset.

And before starting, make a plan with bullet points...



Recommended reading:

**Chain-of-Thought Prompting Elicits Reasoning in Large Language Models.** Wei et al.

**Self-Consistency Improves Chain of Thought Reasoning in Language Models.** Wang et al.

# Planning

Do a linear regression in NumPy in this dataset.

And before starting, make a plan with bullet points...



Agent makes a plan:

Recommended reading:

**Chain-of-Thought Prompting Elicits Reasoning in Large Language Models.** Wei et al.

**Self-Consistency Improves Chain of Thought Reasoning in Language Models.** Wang et al.

# Planning

Do a linear regression in NumPy in this dataset.

And before starting, make a plan with bullet points...



Agent makes a plan:

1. Download and process the dataset
2. Find and understand useful features
3. Write a draft of the code and run it with Python
4. ....

Recommended reading:

**Chain-of-Thought Prompting Elicits Reasoning in Large Language Models.** Wei et al.

**Self-Consistency Improves Chain of Thought Reasoning in Language Models.** Wang et al.

# Planning

Do a linear regression in NumPy in this dataset.

And before starting, make a plan with bullet points...



Agent makes a plan:

1. Download and process the dataset
2. Find and understand useful features
3. Write a draft of the code and run it with Python
4. ....

And then the agent executes the plan

Recommended reading:

**Chain-of-Thought Prompting Elicits Reasoning in Large Language Models.** Wei et al.

**Self-Consistency Improves Chain of Thought Reasoning in Language Models.** Wang et al.

# Planning: Why does it work?

1. Download and process the dataset
2. Find and understand useful features
3. Write a draft of the code and run it with Python
4. ....




Recommended reading:

**Chain-of-Thought Prompting Elicits Reasoning in Large Language Models.** Wei et al.

**Self-Consistency Improves Chain of Thought Reasoning in Language Models.** Wang et al.

# Reflection

Do a linear regression in NumPy in this dataset.



Recommended reading:


**Reflexion: Language Agents with Verbal Reinforcement Learning.** Shinn et al.

# Reflection

Do a linear regression in NumPy in this dataset.



Write



Recommended reading:

**Reflexion: Language Agents with Verbal Reinforcement Learning.** Shinn et al.




# Reflection

Do a linear regression in NumPy in this dataset.



Write

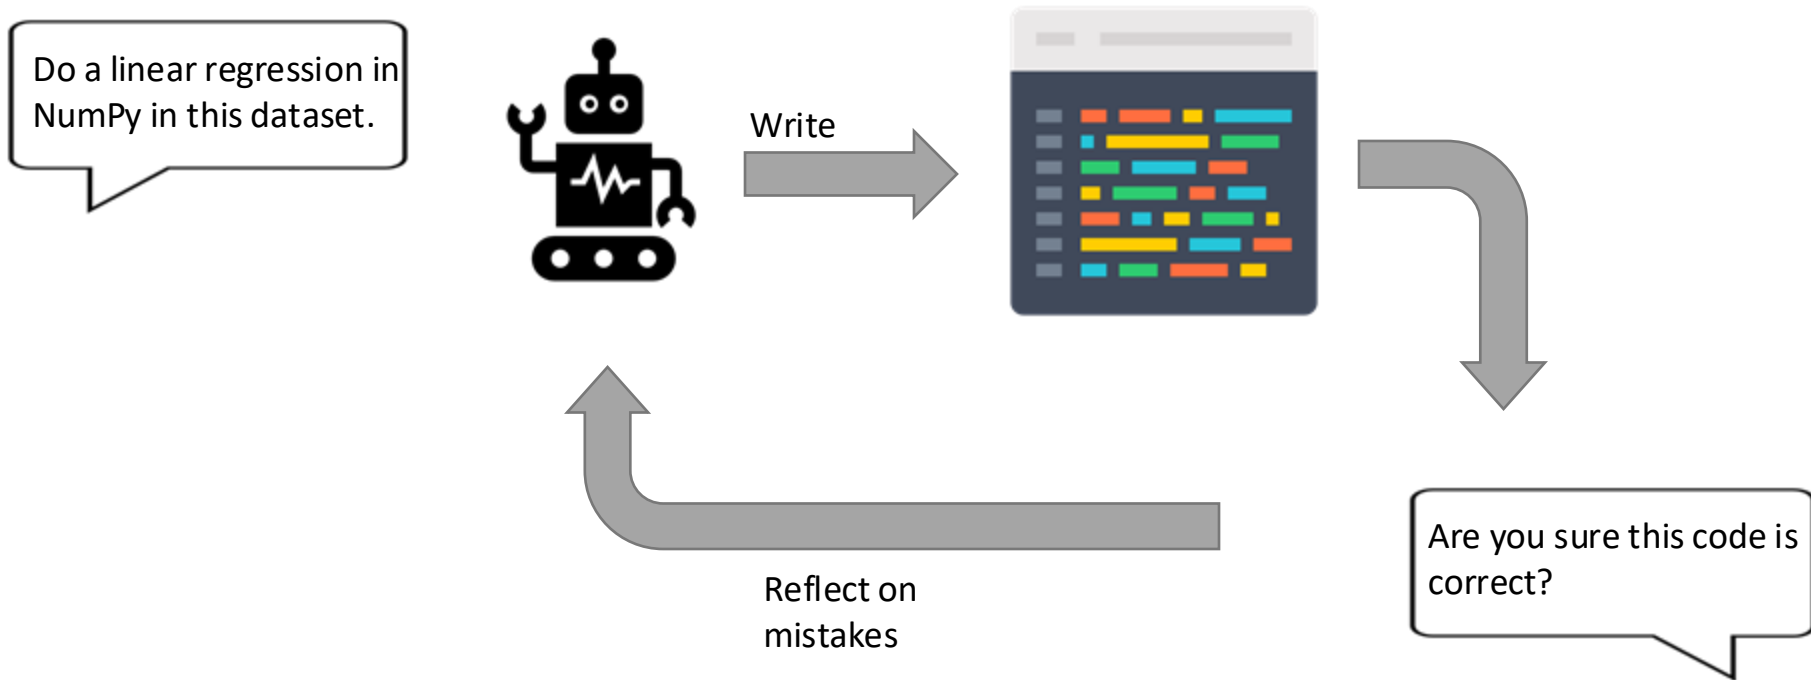


Are you sure this code is correct?

Recommended reading:

**Reflexion: Language Agents with Verbal Reinforcement Learning.** Shinn et al.

# Reflection




Recommended reading:

**Reflexion: Language Agents with Verbal Reinforcement Learning.** Shinn et al.

# Memory

Do a linear regression in NumPy in this dataset.



Recommended reading:

**Enhancing Large Language Models with Long-Term Memory.** Zhong et al.

# Memory

Do a linear regression in NumPy in this dataset.



Previous AI code,  
examples, human code



Recommended reading:


**Enhancing Large Language Models with Long-Term Memory.** Zhong et al.

# Memory

Do a linear regression in NumPy in this dataset.



Has anything  
similar  
happened  
before?



Previous AI code,  
examples, human code




Recommended reading:

**Enhancing Large Language Models with Long-Term Memory.** Zhong et al.

# Memory

Do a linear regression in NumPy in this dataset.



Previous AI code,  
examples, human code




Recommended reading:

**Enhancing Large Language Models with Long-Term Memory.** Zhong et al.

# Memory

Do a linear regression in NumPy in this dataset.



Previous AI code,  
examples, human code



Write the final code

Recommended reading:

**Enhancing Large Language Models with Long-Term Memory.** Zhong et al.

# Multi-Agent

Do a linear regression in NumPy in this dataset.



**Coder  
Agent**



**Critic  
Agent**

Recommended reading:

**ChatDev: Communicative Agents for Software Development.** Qian et al.



# Multi-Agent

Do a linear regression in NumPy in this dataset.



**Coder  
Agent**

Coder Writes the code:



**Critic  
Agent**

Recommended reading:

**ChatDev: Communicative Agents for Software Development.** Qian et al.

# Multi-Agent

Do a linear regression in NumPy in this dataset.



**Coder  
Agent**



Is what the **Coder Agent** did correct?




**Critic  
Agent**

Recommended reading:

**ChatDev: Communicative Agents for Software Development.** Qian et al.

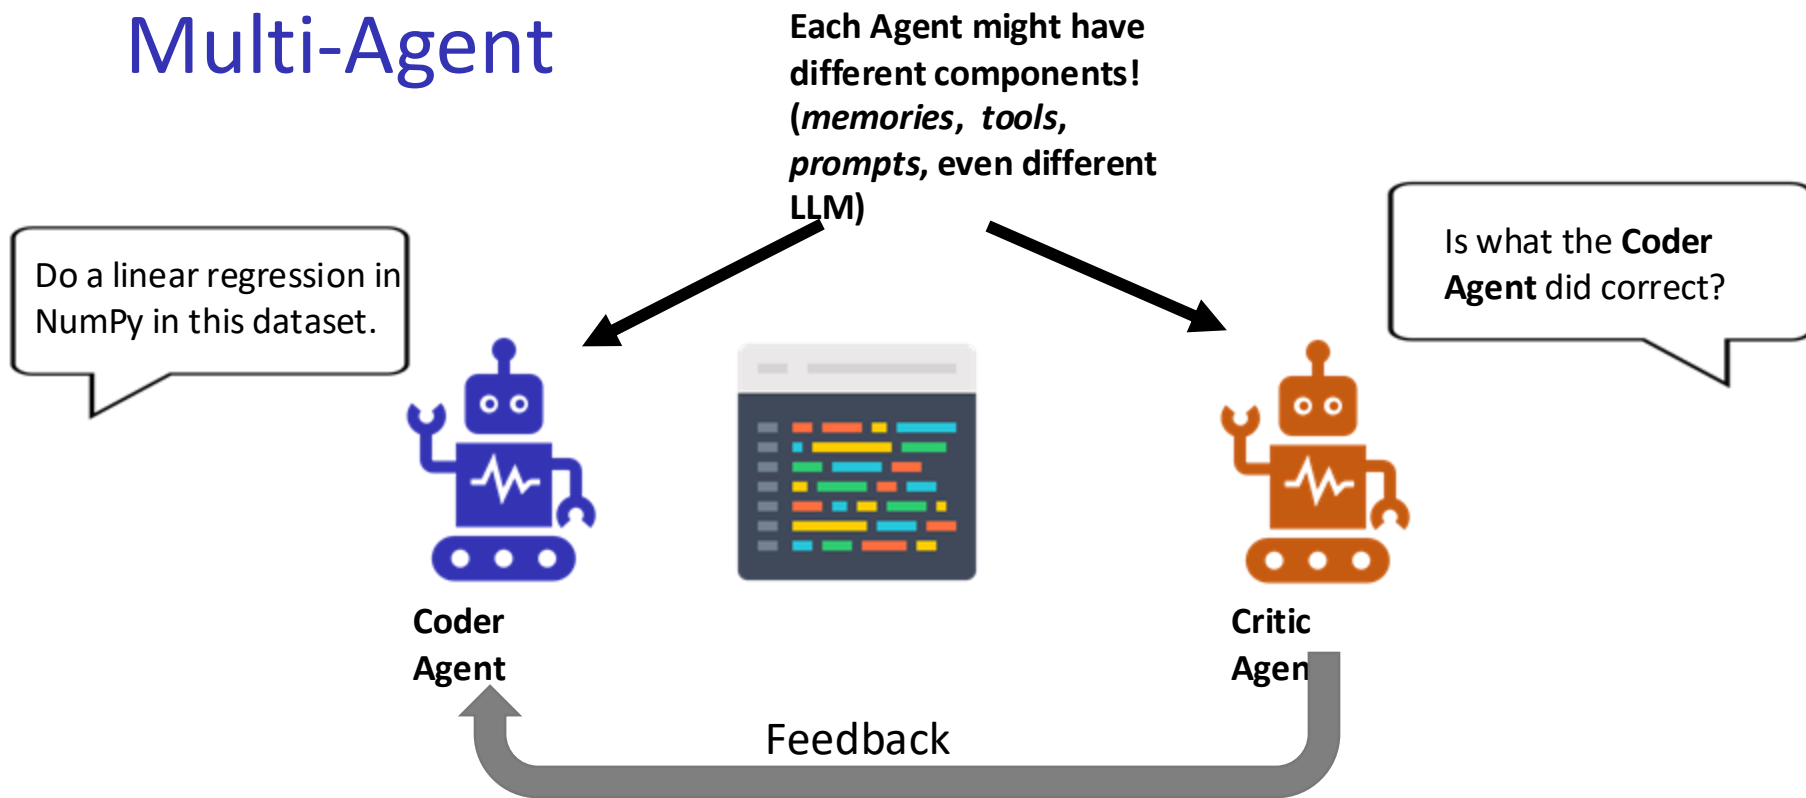
# Multi-Agent



Recommended reading:

**ChatDev: Communicative Agents for Software Development.** Qian et al.

# Multi-Agent



Recommended reading:

**ChatDev: Communicative Agents for Software Development.** Qian et al.

How do we put all these together to create  
powerful AI Agents?

# Pangu-Agent: A Fine-Tunable Generalist Agent with Structured Reasoning

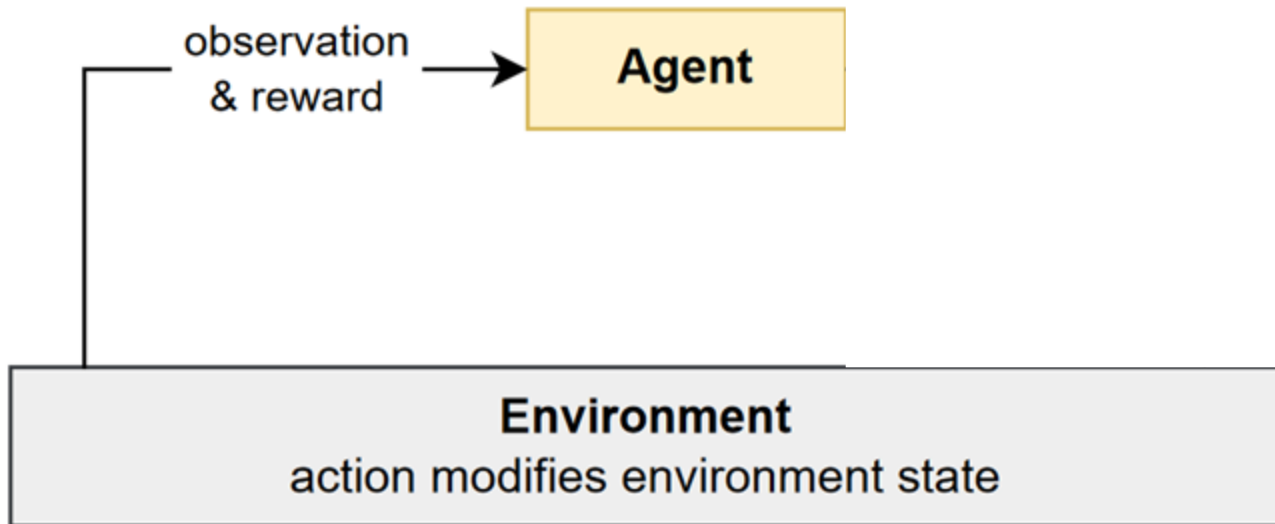
*Filippos Christianos, Georgios Papoudakis, Matthieu Zimmer, Thomas Coste, Zhihao Wu, Jingxuan Chen, Khyati Khandelwal, James Doran, Xidong Feng, Jiacheng Liu, Zheng Xiong, Yicheng Luo, Jianye Hao, Kun Shao, Haitham Bou-Ammar, Jun Wang*



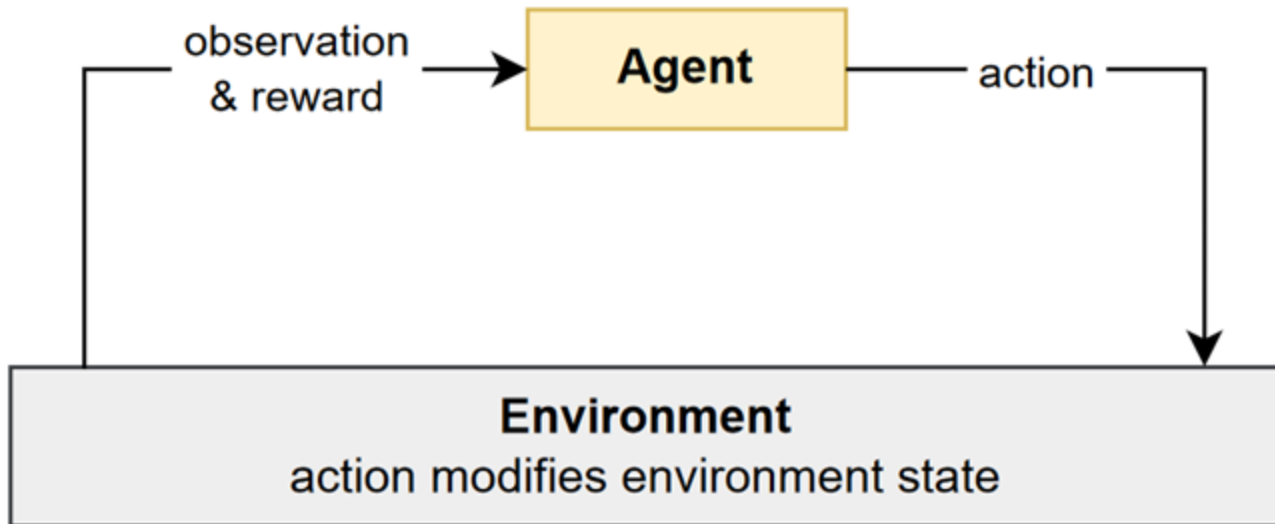
**Agent**


**Environment**

action modifies environment state









What is our objective?

$$\operatorname{argmax}_{\pi} \mathbb{E}[G|\pi]$$

$$\operatorname{argmax}_{\pi} \mathbb{E}[G|\pi]$$

- 1. Solutions are too task specific requiring lots of engineering**

$$\operatorname{argmax}_{\pi} \mathbb{E}[G|\pi]$$

- 1. Solutions are too task specific requiring lots of engineering**
- 2. Agent structures are pre-defined and can't adapt their reasoning to suit tasks**

Here's the state: s

Just give me the answer!

$\pi$

Here's the state:  $s$

Just give me the answer!

$\pi$  ? ?

$$\operatorname{argmax}_{\pi} \mathbb{E}[G|\pi] \longrightarrow$$




$$\operatorname{argmax}_{\pi} \mathbb{E}[G|\pi] \longrightarrow \operatorname{argmax}_{\pi, \mu} \mathbb{E}[G|\pi, \mu]$$

$$\operatorname{argmax}_{\pi} \mathbb{E}[G|\pi] \longrightarrow \operatorname{argmax}_{\pi, \mu} \mathbb{E}[G|\pi, \mu]$$

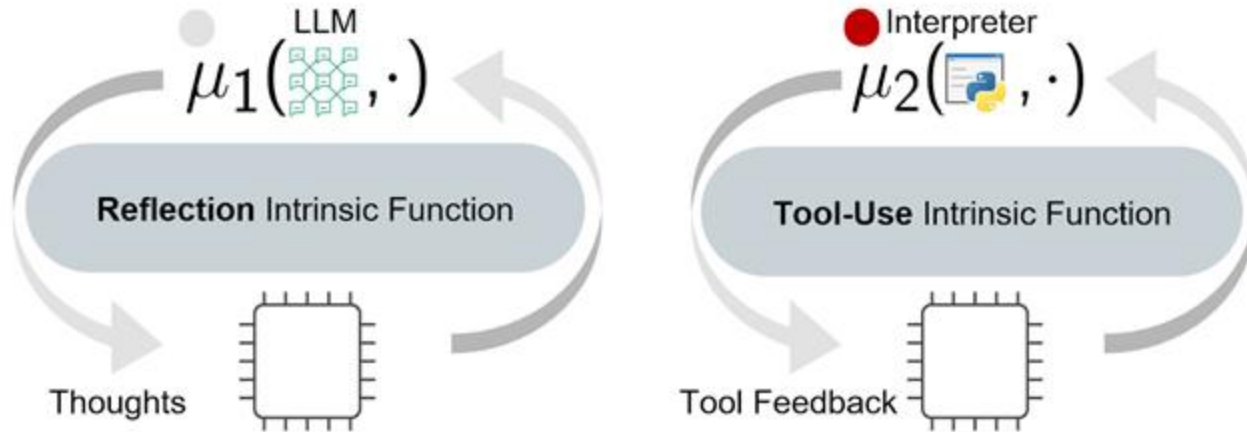
**Structured reasoning:  
intrinsic functions!**

- *Intrinsic function* ( $\mu$ ): one that operates on memory (short term or long term)

- *Intrinsic function*: one that operates on memory (short term or long term)

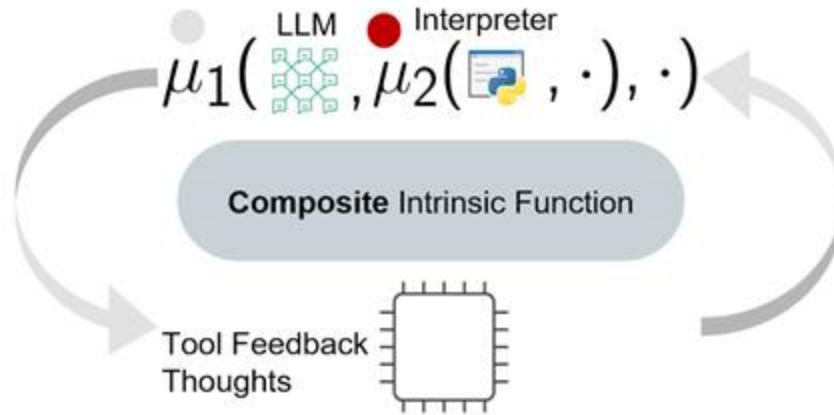


- *Intrinsic function*: one that operates on memory (short term or long term)

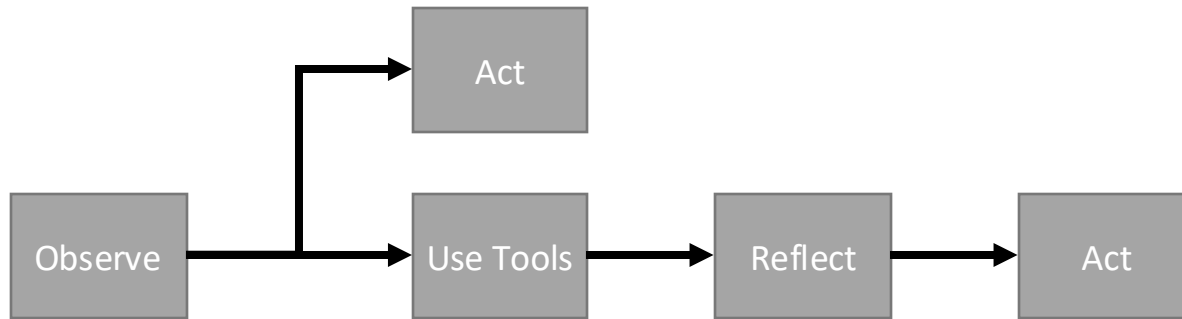


- *Intrinsic function*: one that operates on memory (short term or long term)
- **We can make *composite functions* by *nesting* intrinsic functions**

- *Intrinsic function*: one that operates on memory (short term or long term)
- We can make *composite functions* by *nesting* intrinsic functions

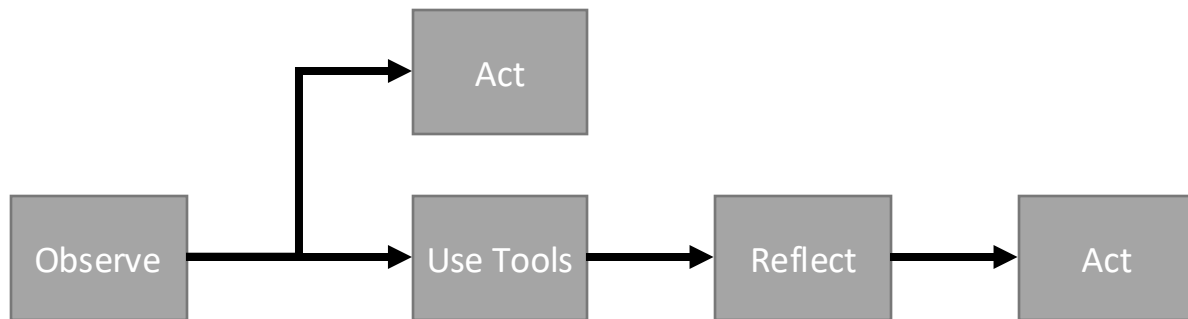


Which means we can make agentic workflows...






Which means we can make agentic workflows...




Still just a composite function!




... And this encompasses everything we have talked about so far...

... And this encompasses everything we have talked about so far...

**Pangu-Agent is built with this formulation in mind**




Results on GSM8K



Results on GSM8K

**And we can implement more complex methods  
(composite) function too..**



Results on GSM8K

# The Pangu Agent framework:



# The Pangu Agent framework:

1. Can use various LLMs. LLMs is just an input to the intrinsic function...

# The Pangu Agent framework:

1. Can use various LLMs. LLMs is just an input to the intrinsic function...
2. It can be used in different environments/tasks.

# The Pangu Agent framework:

1. Can use various LLMs. LLMs is just an input to the intrinsic function...
2. It can be used in different environments/tasks.
3. It can be used with a variety of methods... (including existing ones)

**But that's not all... Remember**  $\operatorname{argmax}_{\pi, \mu} \mathbb{E}[G | \pi, \mu]$

# ALFWorld



# Let's look at this pipeline ...

Our Agent



Nested-Order Interaction

$$\pi\left(\mu.(\cdot)\right)$$

Think

$\mu.(\cdot)$

Reflect


Other orders of interaction ...

Interaction

$$\pi(\mu_1(\mu_0(\mathbf{o}_t, \mathbf{a}_t), \text{arg}))$$


Repeat

Mix-n-Match Nested Interactions



Mix-n-Match

Supervised Fine-Tune  
(Next token pred)



Mixed-data set



Training Tasks

Tasks	Openchat v3.5				
	Direct	FS-CoT	1-step SFT	2-step SFT	3-step SFT
<b>ALFWorld</b>	0.04	0.22	0.45	0.68	<b>0.82</b>

Table 4: Benchmark of Openchat v3.5 with/without fine-tuning on held-out tasks.

# And then RL...

Tasks	Openchat v3.5		Llama-2-7B			
	Direct	FS-CoT	Original	SFT	SFT+RL	RL
<b>ALFWorld</b>	0.04	0.22	0	0.5	<b>0.88</b>	0.04
<b>BabyAI-GoToObj-v0</b>	0.31	0.61	0.28	0.75	<b>0.91</b>	0.87
<b>BabyAI-GoToRedBlueBall-v0</b>	0.11	0.43	0.04	0.21	<b>0.77</b>	0.69

Table 5: Benchmark of Openchat and LLama-2-7b with/without fine-tuning on held-out tasks.



Openchat SFT

Task: Put clean mug  
in coffeemachine

Press **q** to Exit

# Part V: App Control with AI Agents

**GOAL:** I'd like to get a new three-seater sofa for Christmas because my old one broke – search in the “industry buy” app.

**GOAL:** I'd like to get a new three-seater sofa for Christmas because my old one broke – search in the “industry buy” app.



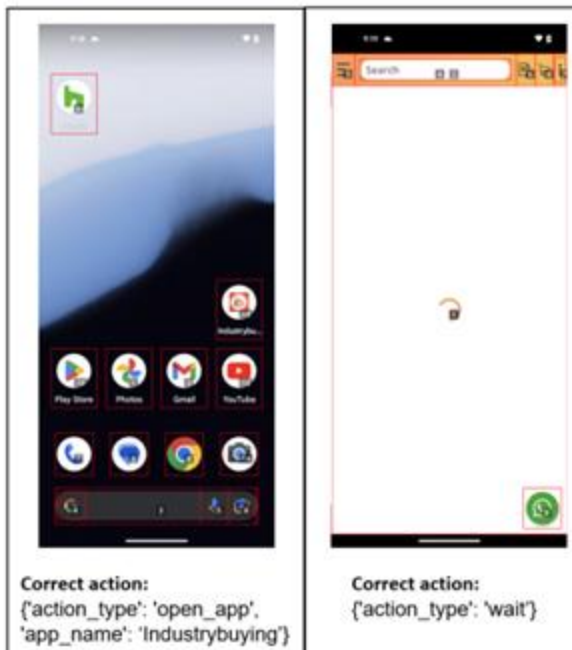
**GOAL:** I'd like to get a new three-seater sofa for Christmas because my old one broke – search in the “industry buy” app.




Correct action:

```
{  
  'action_type': 'open_app',  
  'app_name': 'Industrybuying'  
}
```






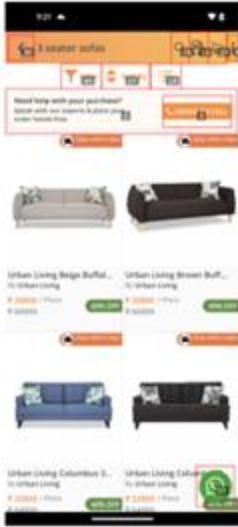
**GOAL:** I'd like to get a new three-seater sofa for Christmas because my old one broke – search in the “industry buy” app.



**GOAL:** I'd like to get a new three-seater sofa for Christmas because my old one broke – search in the “industry buy” app.



**GOAL:** I'd like to get a new three-seater sofa for Christmas because my old one broke – search in the “industry buy” app.

					
Correct action: { <code>'action_type': 'open_app',</code> <code>'app_name': 'Industrybuying'</code> }	Correct action: { <code>'action_type': 'wait'</code> }	Correct action: { <code>'action_type': 'click',</code> <code>'target_element': 27</code> }	Correct action: { <code>'action_type': 'input_text',</code> <code>'text': '3 seater sofas'</code> }	Correct action: { <code>'action_type': 'click',</code> <code>'target_element': 24</code> }	End of episode



# Training Datasets

1. Android in the Wild (AitW): 17k\* episodes
2. Android Control: 15k episodes

# Training Datasets

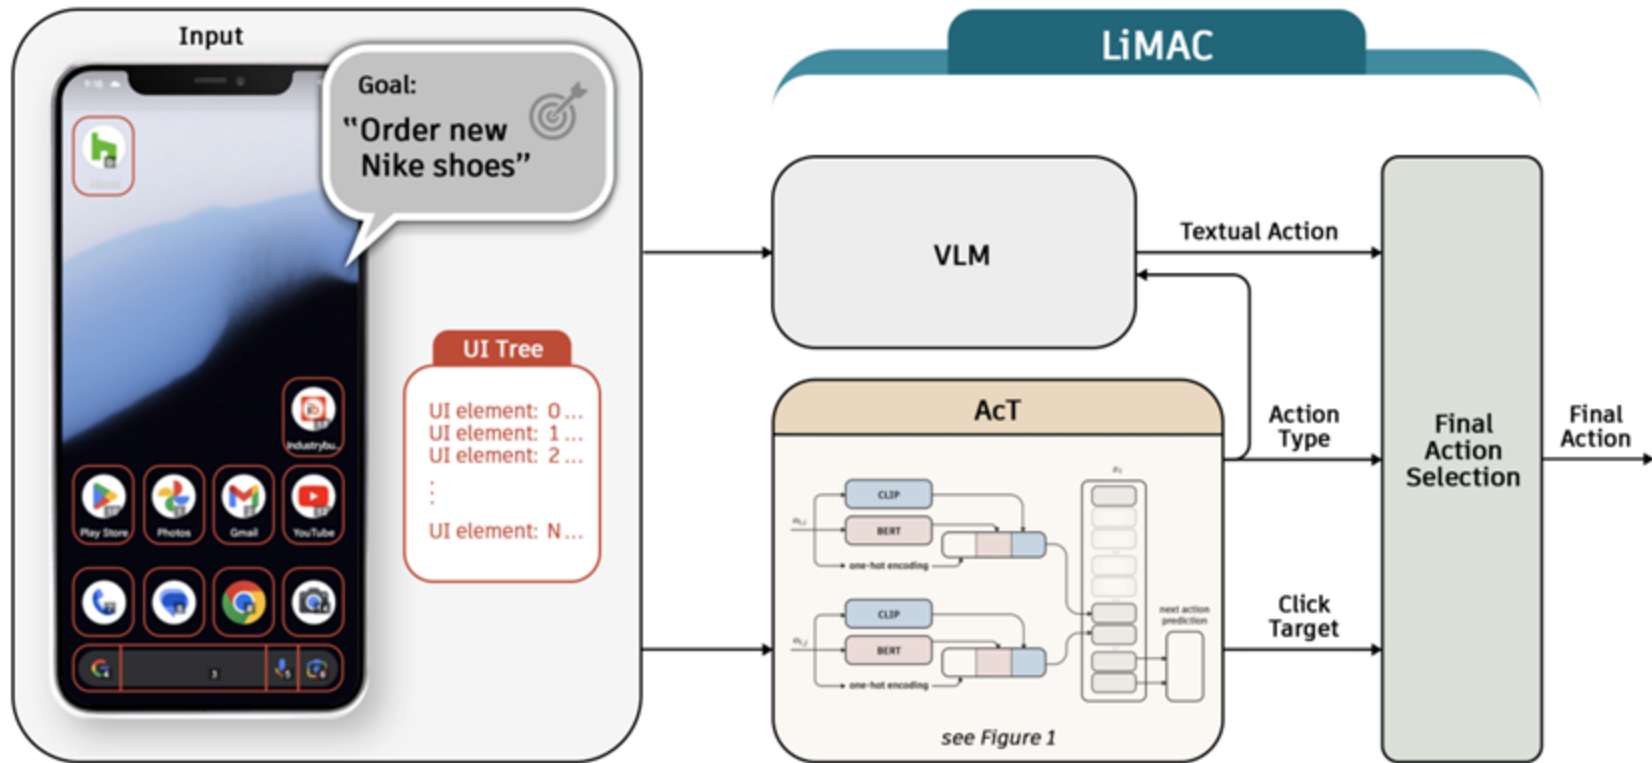
1. Android in the Wild (AitW): 17k\* episodes
2. Android Control: 15k episodes

Each episode contains:

- Goal
- Observations (screenshot, or UI tree)
- Actions (click, type, wait, etc...)

# Motivation

1. Low Computational requirements
2. Efficient preprocessing of input  
screenshot
3. Condition action in past trajectory
4. Specialisation towards click actions



# Approach

1. Crop each UI element and embed it into one token
  - each observation becomes a list of tokens

# Approach

1. Crop each UI element and embed it into one token
  - each observation becomes a list of tokens
2. Formulate the action type selection as a classification problem
  - train a small, fast, transformer called AcT

# Approach

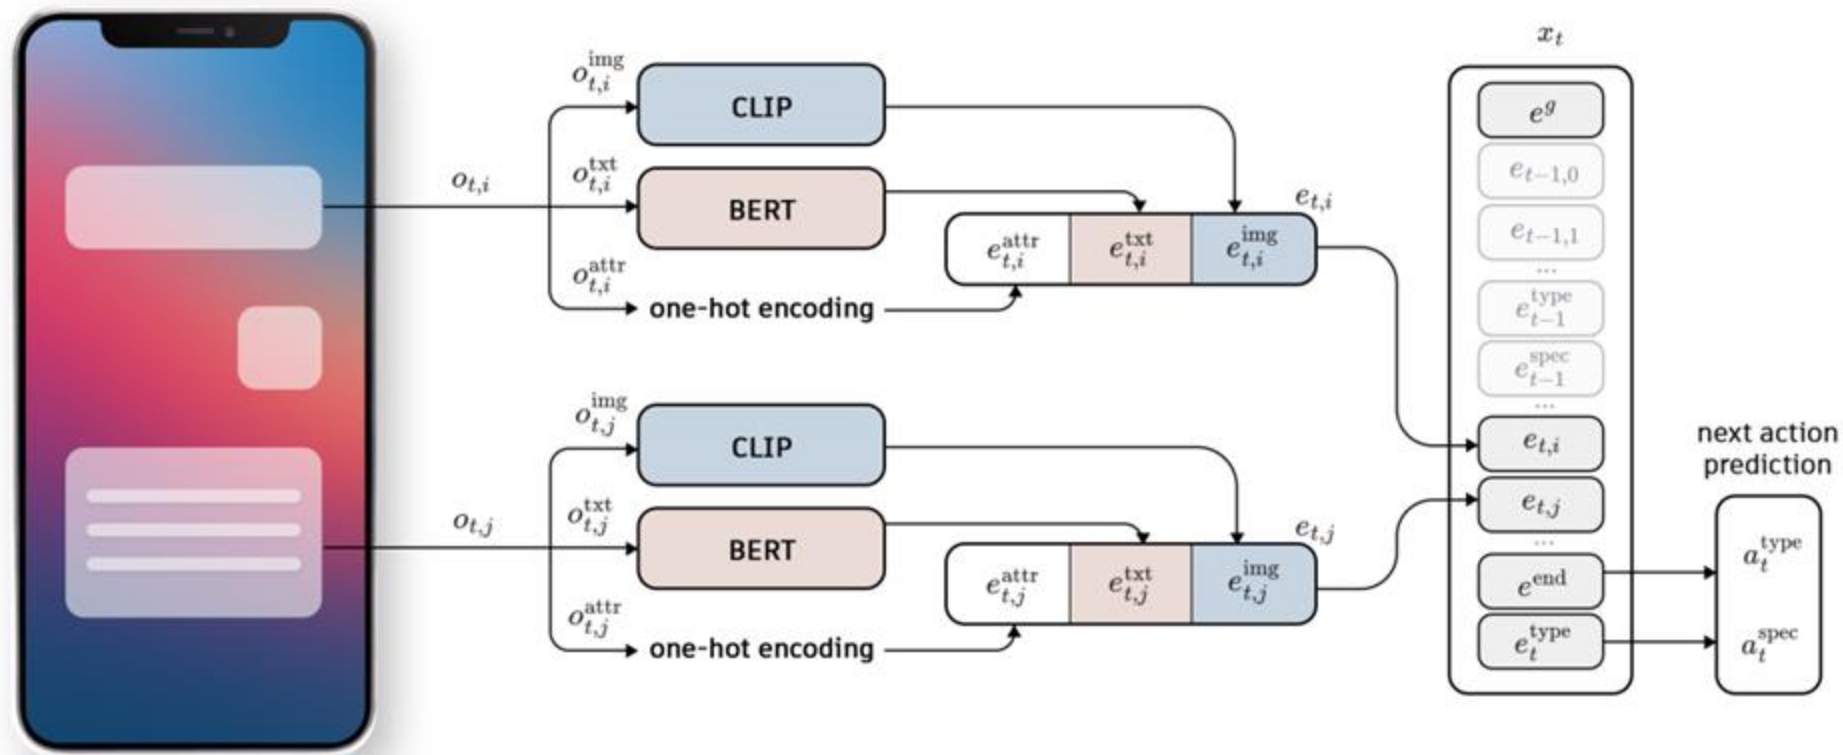
1. Crop each UI element and embed it into one token
  - each observation becomes a list of tokens
2. Formulate the action type selection as a classification problem
  - train a small, fast, transformer called AcT
3. Use the same model with contrastive learning to predict the target of click actions

# Approach

1. Crop each UI element and embed it into one token
  - each observation becomes a list of tokens
2. Formulate the action type selection as a classification problem
  - train a small, fast, transformer called AcT
3. Use the same model with contrastive learning to predict the target of click actions
4. Use a small VLM for actions that require text



phone screen at timestep  $t$



# Action Transformer Training

The action type prediction head over the 10 possible action types. It is trained using to maximise the log likelihood of the correct action type

$$\mathcal{L}_{\text{type}} = -\mathbf{E}_{a^{\text{type}}, x \in \mathcal{D}} [\log(p(a^{\text{type}}|h))]$$

# Action Transformer Training

We also project the hidden states of AcT to a space that we will perform contrastive learning. Compare the hidden state of the action type embedding with the hidden state of all UI elements.

$$q^{\text{type}} = f_{\text{target}}(h_t^{\text{type}}) \quad \text{and} \quad p^{\text{ui}} = f_{\text{target}}(h^{\text{ui}}) \quad S = \frac{qp^T}{\|q\| \cdot \|p\|_r} \tau$$

# Action Transformer Training

We also project the hidden states of AcT to a space that we will perform contrastive learning. Compare the hidden state of the action type embedding with the hidden state of all UI elements.

$$q^{\text{type}} = f_{\text{target}}(h_t^{\text{type}}) \quad \text{and} \quad p^{\text{ui}} = f_{\text{target}}(h^{\text{ui}}) \quad S = \frac{qp^T}{\|q\| \cdot \|p\|_r} \tau$$

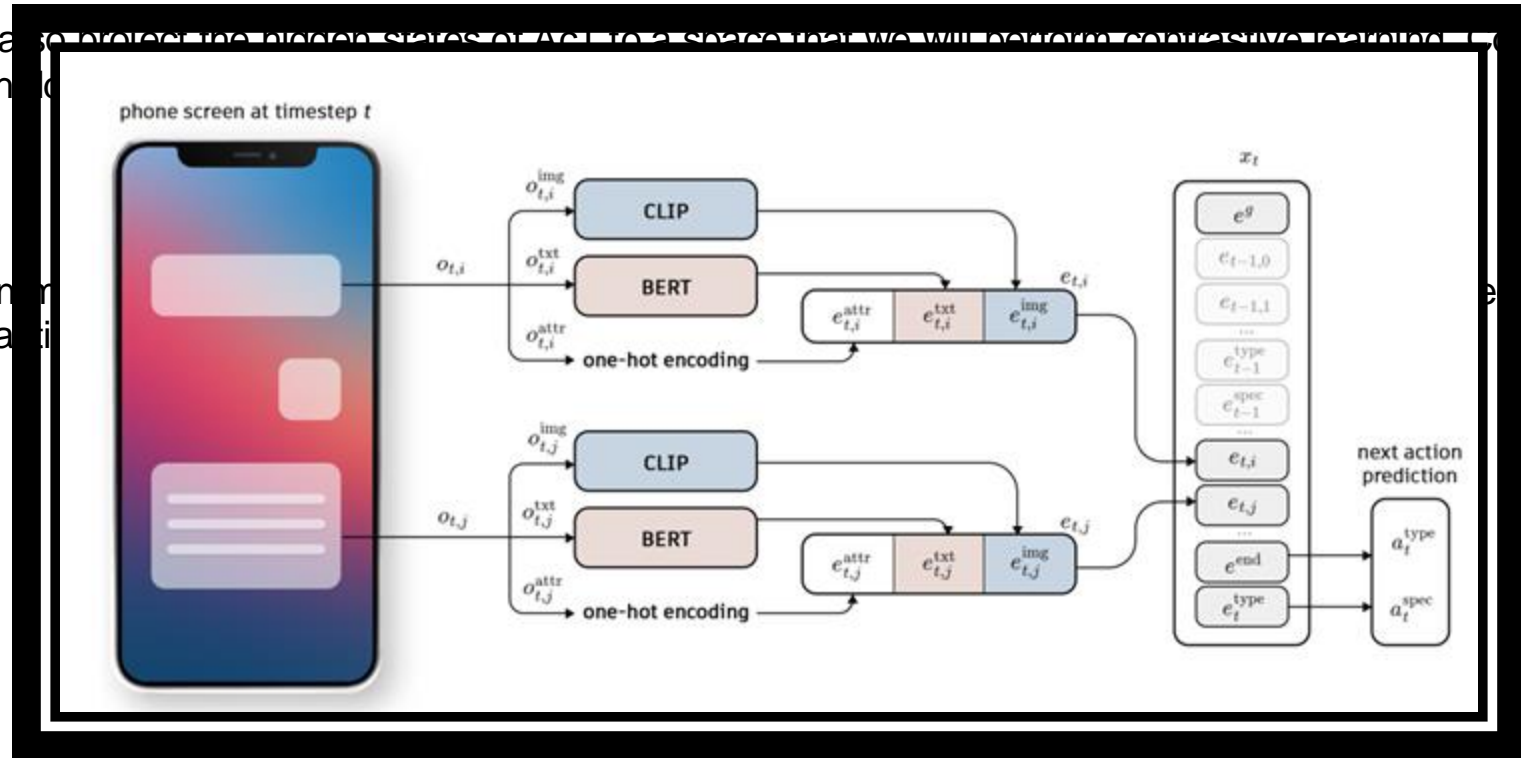
Then minimise the InfoNCE loss to project the correct target element embedding close to embedding of the action type

$$\mathcal{L}_{\text{elem}} = -\mathbf{E} \left[ \log \frac{\exp(S_+)}{\sum_{i=1}^K \exp(S_i)} \right]$$

# Action Transformer Training

We also process the action series of the screen. We will extract consecutive features to compare the history.

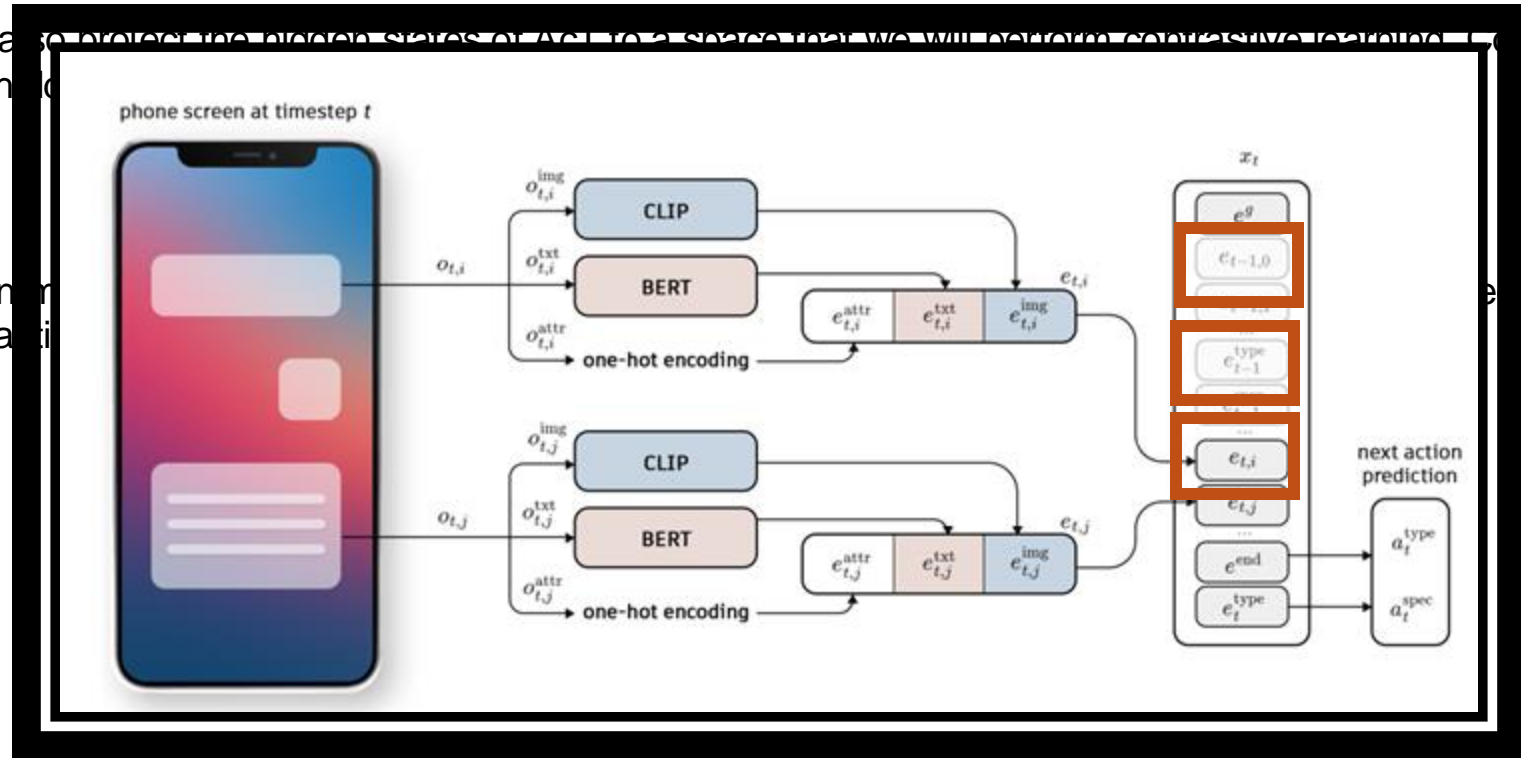
Then, we will process the action series of the screen. We will extract consecutive features to compare the action.



# Action Transformer Training

We also process the action series of the screen before we will perform contrastive learning to compare the hidden states.

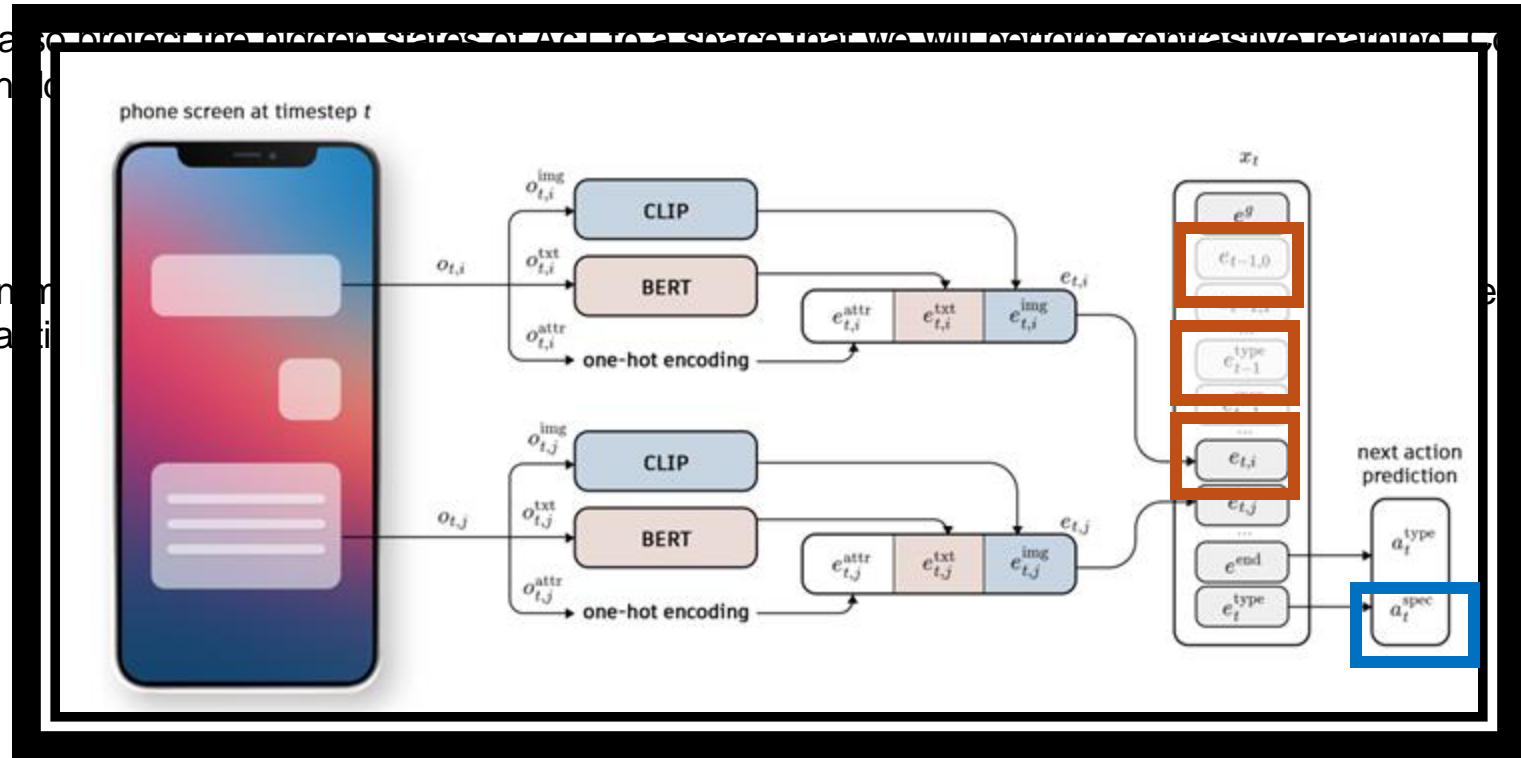
Then we process the action series of the screen before we will perform contrastive learning to compare the hidden states.



# Action Transformer Training

We also process the action series of the screen before we will perform contrastive learning. We compare the hidden states of the action series with the hidden states of the screen series.

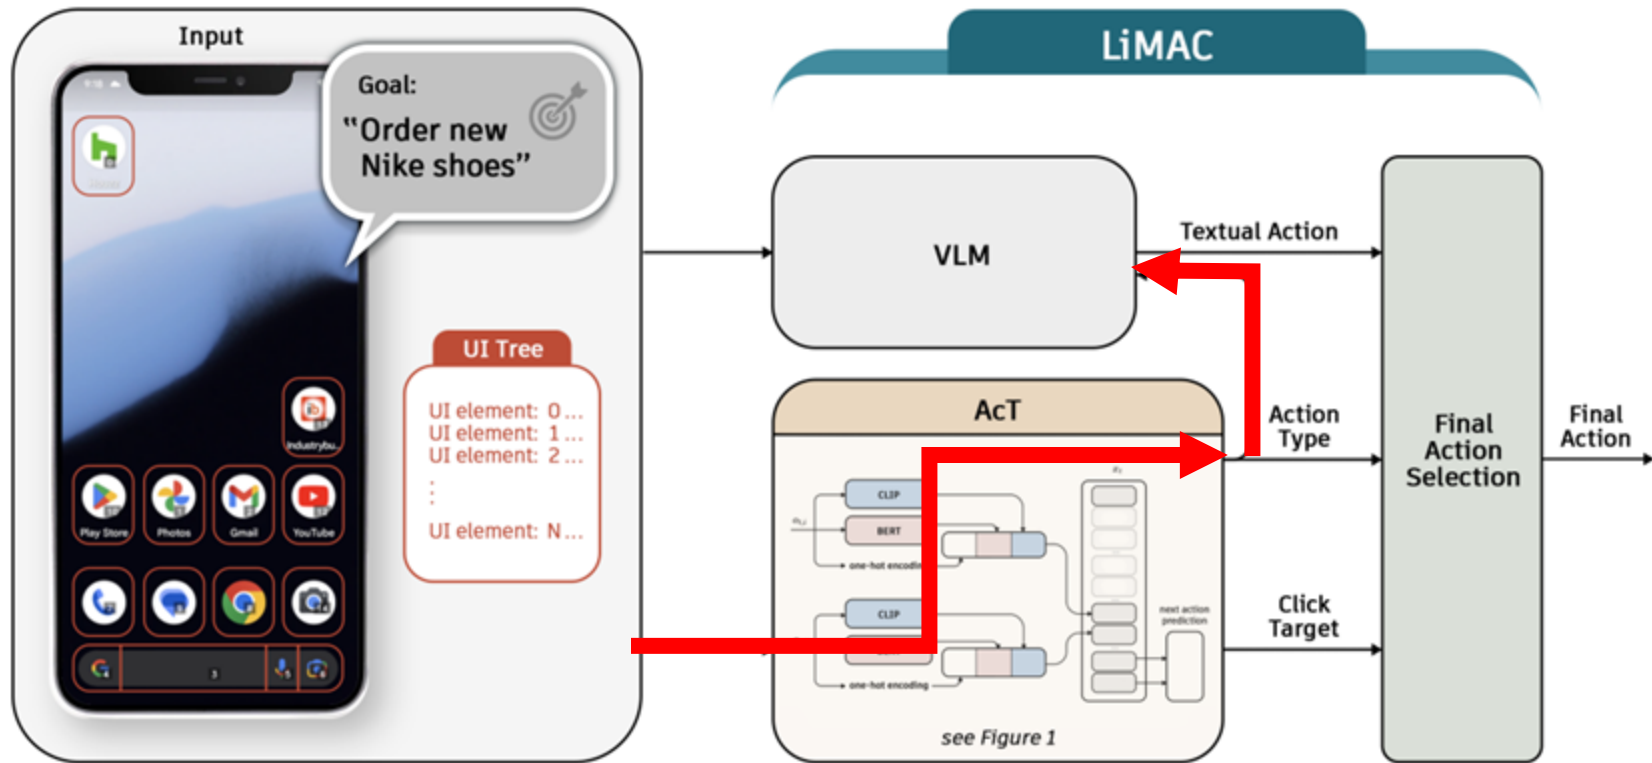
Then we process the action series of the screen before we will perform contrastive learning. We compare the hidden states of the action series with the hidden states of the screen series.



# VLM fine-tuning

- LiMAC uses AcT to predict the action type and the target element of click actions.
- Actions that require text for the specifications are handled by a small VLM.
- The VLM receives as input the screenshot with bounding boxes around the UI elements with a corresponding number as well as the goal in natural language
- The VLM is fine-tuned using the training dataset in all actions
- The VLM is fine-tuned to maximise the log likelihood
- We fine-tune two different VLMs:
  - Florence2 which is 820M parameters
  - Qwen2-VL which is 2B parameters





# Overall Accuracy Evaluation

Model	Size ↓	Avg Inf. Time (s) ↓	Overall ↑	
			AitW	AndCtrl
SeeAct <sub>choice</sub>	unk	9.81	37.7	29.9
SeeAct <sub>ann</sub>	unk	9.76	42.5	35.5
T3A	unk	4.87	26.9	53.1
M3A	unk	10.64	35.6	57.5
Florence2	820M	0.50	70.8	57.0
LiMAC with Florence2 (ours)	+520M	<b>0.34</b>	<b>72.2</b>	<b>63.1</b>
Qwen2-VL	2B	3.03	51.0	52.2
LiMAC with Qwen2-VL (ours)	+520M	0.63	70.9	62.5

# Overall Accuracy Evaluation

Model	Size ↓	Avg Inf. Time (s) ↓	Overall ↑	
			AitW	AndCtrl
SeeAct <sub>choice</sub>	unk	9.81	37.7	29.9
SeeAct <sub>ann</sub>	unk	9.76	42.5	35.5
T3A	unk	4.87	26.9	53.1
M3A	unk	10.64	35.6	57.5
Florence2	820M	0.50	70.8	57.0
LiMAC with Florence2 (ours)	+520M	<b>0.34</b>	<b>72.2</b>	<b>63.1</b>
Qwen2-VL	2B	3.03	51.0	52.2
LiMAC with Qwen2-VL (ours)	+520M	0.63	70.9	62.5

# Ablation Study (Action-Type and Click-Target Accuracy)







Framework	Modules Used			Action Type		Click Target		Text	
	Type	Click	Text	AitW	AndCtrl	AitW	AndCtrl	AitW	AndCtrl
SeeAct only	SeeAct <sub>choice</sub>	SeeAct <sub>choice</sub>	SeeAct <sub>choice</sub>	67.1	66.8	36.9	48.5	69.4	67.1
SeeAct only	SeeAct <sub>ann</sub>	SeeAct <sub>ann</sub>	SeeAct <sub>ann</sub>	68.2	66.8	44.7	55.7	66.0	61.8
T3A only	T3A	T3A	T3A	56.2	67.7	33.5	71.1	66.5	<b>78.4</b>
M3A only	M3A	M3A	M3A	63.8	69.8	48.3	<b>77.1</b>	67.3	74.3
Qwen only	Qwen2-VL	Qwen2-VL	Qwen2-VL	81.7	70.7	53.2	55.2	70.5	75.7
LiMAC (ours)	AcT	Qwen2-VL	Qwen2-VL	<b>86.9</b>	<b>82.3</b>	53.2	55.2	70.5	75.7
LiMAC (ours)	AcT	AcT	Qwen2-VL	<b>86.9</b>	<b>82.3</b>	<b>77.4</b>	65.4	70.5	75.7
Florence only	Florence2	Florence2	Florence2	86.4	79.6	76.2	62.0	<b>84.2</b>	77.5
LiMAC (ours)	AcT	Florence2	Florence2	<b>86.9</b>	<b>82.3</b>	76.2	62.0	<b>84.2</b>	77.5
LiMAC (ours)	AcT	AcT	Florence2	<b>86.9</b>	<b>82.3</b>	<b>77.4</b>	65.4	<b>84.2</b>	77.5

# Ablation Study (Action-Type and Click-Target Accuracy)

Framework	Modules Used			Action Type		Click Target		Text	
	Type	Click	Text	AitW	AndCtrl	AitW	AndCtrl	AitW	AndCtrl
SeeAct only	SeeAct <sub>choice</sub>	SeeAct <sub>choice</sub>	SeeAct <sub>choice</sub>	67.1	66.8	36.9	48.5	69.4	67.1
SeeAct only	SeeAct <sub>ann</sub>	SeeAct <sub>ann</sub>	SeeAct <sub>ann</sub>	68.2	66.8	44.7	55.7	66.0	61.8
T3A only	T3A	T3A	T3A	56.2	67.7	33.5	71.1	66.5	<b>78.4</b>
M3A only	M3A	M3A	M3A	63.8	69.8	48.3	<b>77.1</b>	67.3	74.3
Qwen only	Qwen2-VL	Qwen2-VL	Qwen2-VL	81.7	70.7	53.2	55.2	70.5	75.7
LiMAC (ours)	AcT	Qwen2-VL	Qwen2-VL	<b>86.9</b>	<b>82.3</b>	53.2	55.2	70.5	75.7
LiMAC (ours)	AcT	AcT	Qwen2-VL	<b>86.9</b>	<b>82.3</b>	<b>77.4</b>	65.4	70.5	75.7
Florence only	Florence2	Florence2	Florence2	86.4	79.6	76.2	62.0	<b>84.2</b>	77.5
LiMAC (ours)	AcT	Florence2	Florence2	<b>86.9</b>	<b>82.3</b>	76.2	62.0	<b>84.2</b>	77.5
LiMAC (ours)	AcT	AcT	Florence2	<b>86.9</b>	<b>82.3</b>	<b>77.4</b>	65.4	<b>84.2</b>	77.5







# Case Study of an Episode

**Goal:** I want to look for a flight from Detroit to Las Vegas in business class for 4 passengers on Expedia departing October 11, 2023 and returning October 16, 2023 because I'm organising a family vacation to Las Vegas.

					
<p><b>Correct action:</b> {<code>'action_type': 'open_app',</code> <code>'app_name': 'Expedia'</code>}</p> <p><b>Predicted action:</b> {<code>'action_type': 'open_app',</code> <code>'app_name': 'Expedia'</code>}</p>	<p><b>Correct action:</b> {<code>'action_type': 'click',</code> <code>'target_element': 17</code>}</p> <p><b>Predicted action:</b> {<code>'action_type': 'click',</code> <code>'target_element': 17</code>}</p>	<p><b>Correct action:</b> {<code>'action_type': 'click',</code> <code>'target_element': 23</code>}</p> <p><b>Predicted action:</b> {<code>'action_type': 'click',</code> <code>'target_element': 23</code>}</p>	<p><b>Correct action:</b> {<code>'action_type': 'input_text',</code> <code>'text': 'Detroit'</code>}</p> <p><b>Predicted action:</b> {<code>'action_type': 'input_text',</code> <code>'text': 'Detroit'</code>}</p>	<p><b>Correct action:</b> {<code>'action_type': 'click',</code> <code>'target_element': 27</code>}</p> <p><b>Predicted action:</b> {<code>'action_type': 'click',</code> <code>'target_element': 27</code>}</p>	<p><b>Correct action:</b> {<code>'action_type': 'input_text',</code> <code>'text': 'Las Vegas'</code>}</p> <p><b>Predicted action:</b> {<code>'action_type': 'input_text',</code> <code>'text': 'Detroit'</code>}</p>


# Case Study of an Episode

**Goal:** I want to look for a flight from Detroit to Las Vegas in business class for 4 passengers on Expedia departing October 11, 2023 and returning October 16, 2023 because I'm organising a family vacation to Las Vegas.

					
<p>Correct action: {<code>'action_type': 'open_app',</code> <code>'app_name': 'Expedia'</code>}</p> <p>Predicted action: {<code>'action_type': 'open_app',</code> <code>'app_name': 'Expedia'</code>}</p>	<p>Correct action: {<code>'action_type': 'click',</code> <code>'target_element': 17</code>}</p> <p>Predicted action: {<code>'action_type': 'click',</code> <code>'target_element': 17</code>}</p>	<p>Correct action: {<code>'action_type': 'click',</code> <code>'target_element': 21</code>}</p> <p>Predicted action: {<code>'action_type': 'click',</code> <code>'target_element': 21</code>}</p>	<p>Correct action: {<code>'action_type': 'input_text',</code> <code>'text': 'Detroit'</code>}</p> <p>Predicted action: {<code>'action_type': 'input_text',</code> <code>'text': 'Detroit'</code>}</p>	<p>Correct action: {<code>'action_type': 'click',</code> <code>'target_element': 27</code>}</p> <p>Predicted action: {<code>'action_type': 'click',</code> <code>'target_element': 27</code>}</p>	<p>Correct action: {<code>'action_type': 'input_text',</code> <code>'text': 'Detroit'</code>}</p> <p>Predicted action: {<code>'action_type': 'input_text',</code> <code>'text': 'Detroit'</code>}</p>



# Multi-Agent RL – Resources



## Multi-Agent Reinforcement Learning: An Introduction

Find it online for free! [www.marl-book.com](http://www.marl-book.com)

**Code repo:** [github.com/uea-agents](https://github.com/uea-agents)

- 25 active code repos
- Extended PyMARL + blog post

**Book codebase:** [github.com/marl-book/codebase](https://github.com/marl-book/codebase)

## MULTI-AGENT REINFORCEMENT LEARNING

FOUNDATIONS AND  
MODERN APPROACHES



Stefano V. Albrecht  
Filippos Christianos  
Lukas Schäfer